

Executive Summary

Parallel Programming

Recitation Session 5

Thomas Weibel <weibelt@ethz.ch>

Laboratory for Software Technology,
Swiss Federal Institute of Technology Zürich

April 1, 2010

- Data partitioning with parallel matrix multiplication
- How to manage threads

Matrix Multiplication

Thomas Weibel <weibelt@ethz.ch>

Parallel Programming
Matrix Multiplication

2

Outline


Parallel Matrix Multiplication

Data partitioning based on

- Output matrix **C**
- Input matrix **A** and input matrix **B**

1 Matrix Multiplication

2 Thread Pools



Source: <http://www.xkcd.com>

Output Partitioning (2 Threads)

```
// Thread 0
for (i=0; i<N/2; i++) {
    for (j=0; j<N; j++) {
        for (k=0; k<N; k++) {
            c[i][j] += a[i][k] * b[k][j];
        }
    }
}
// Thread 1
for (i=N/2; i<N; i++) {
    for (j=0; j<N; j++) {
        for (k=0; k<N; k++) {
            c[i][j] += a[i][k] * b[k][j];
        }
    }
}
```

Input Partitioning: Error in Last Week's Slides

```
// Thread 0
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        for (k=0; k<N/2; k++) {
            c[i][j] += a[i][k] * b[k][j];
        }
    }
}
// Thread 1
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        for (k=N/2; k<N; k++) {
            c[i][j] += a[i][k] * b[k][j];
        }
    }
}
```

Error: Possible race condition when writing `c[i][j]`

Input Partitioning: Locking

```
Object[][] lock; // Lock "matrix"
// Thread 0
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        synchronized(lock[i][j]) {
            for (k=0; k<N/2; k++) {
                c[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
// Thread 1
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        synchronized(lock[i][j]) {
            for (k=N/2; k<N; k++) {
                c[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
```

Overhead?

- A complete row is locked
- Actual lock contention will be moderate to low
- In practice the slow-down – with respect to output partitioning – is moderate (only a few percent)

Input Partitioning: Fine-Grain Locking

Outline

```
Object[][] lock; // Lock "matrix"
// Thread 0
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        for (k=0; k<N/2; k++) {
            synchronized(lock[i][j]) {
                c[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
// Thread 1
for (i=0; i<N; i++) {
    for (j=0; j<N; j++) {
        for (k=N/2; k<N; k++) {
            synchronized(lock[i][j]) {
                c[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
```

Significant overhead:
About 3 times slower than
coarse-grain locking


1 Matrix Multiplication

2 Thread Pools

The Art of Multiprocessor Programming

Thread Overhead

Source of the following material about Thread Pools: "The Art of Multiprocessor Programming" by Maurice Herlihy and Nir Shavit



- Threads require resources
 - Memory for stacks
 - Setup, teardown
- Scheduler overhead
- Worse for short-lived threads

Thread Pools

Thread Pool = Abstraction

- More sensible to keep a pool of long-lived threads
- Threads assigned short-lived tasks
 - Runs the task
 - Rejoins pool
 - Waits for next assignment
- Insulate programmer from platform
 - Big machine, big pool
 - And vice-versa
- Portable code
 - Runs well on any platform
 - No need to mix algorithm/platform concerns

ExecutorService Interface

Future<T>

In `java.util.concurrent`:

- Task = Runnable object
 - If no result value expected
 - Calls `run()` method.
- Task = Callable<T> object
 - If result value of type T expected
 - Calls `T call()` method.

```
Callable<T> task = ...;
...
Future<T> future = executor.submit(task);
...
T value = future.get();
```

- Submitting a Callable<T> task returns a Future<T> object
- The Future's `get()` method blocks until the value is available

Future<?>

Note

```

Runnable task = ...;
...
Future<?> future = executor.submit(task);
...
future.get();

```

- Submitting a Runnable task returns a Future<?> object
- The Future's get() method blocks until the computation is complete

- Executor Service submissions are purely advisory in nature
- The executor
 - Is free to ignore any such advice
 - And could execute tasks sequentially ...

Fibonacci

Disclaimer

$$F(n) := \begin{cases} 1, & n = 0 \\ 1, & n = 1 \\ F(n-1) + F(n-2), & n > 1 \end{cases}$$

- Potential parallelism
- Dependencies

- This Fibonacci implementation is very inefficient
 - So don't deploy it!
- But illustrates our point
 - How to deal with dependencies

Multithreaded Fibonacci

Summary

```
class FibTask implements Callable<Integer> {
    static ExecutorService exec =
        Executors.newCachedThreadPool();
    int arg;
    public FibTask(int n) {
        arg = n;
    }
    public Integer call() {
        if (arg > 2) {
            Future<Integer> left =
                exec.submit(new FibTask(arg-1));
            Future<Integer> right =
                exec.submit(new FibTask(arg-2));
            return left.get() + right.get();
        } else {
            return 1;
        }
    }
}
```

Enjoy your vacation!

